

AI Builder Take-Home

What this is (30-second read)

At Scaler, sales happens on a phone call. A BDA warms up a lead, calls them, tries to understand their goals, answers their doubts, and nudges them to take an entrance test. Most drop off — before the call, during the call, or after. Each drop-off is a place where AI can supercharge the BDA.

Your job: build an agent that supercharges two of those moments — the pre-sales WhatsApp nudge to the BDA (warming them up about the lead before the call, driving pickup and call quality), and the post-call personalised PDF to the lead (building enough trust to get them to the test). Both delivered on WhatsApp.

5 hours of work. We respect your time and we'll read it as a judgment signal if you go significantly over.

Why we're hiring for this

Scaler is going AI-native end-to-end. Sales is where we're starting. If your prototype makes one BDA meaningfully more effective at converting leads into enrolments, we roll it out and you own the roadmap. This is a real slice of that roadmap, not a puzzle.

An invitation, up front

If you read this brief and have a genuinely better idea for how AI should supercharge a BDA at Scaler — build that instead. Tell us in your README what you did differently and why. The strongest submissions usually aren't the ones that follow the brief exactly. We want builders with taste, not executors.

How sales at Scaler works

1. **Lead is created** — through a marketing campaign (MC), organic search, or re-activation of an old lead that went cold

2. **Lead is assigned to a BDA** — a Business Development Associate on the sales team
 3. **Pre-call** — today, the BDA usually walks into the call with nothing more than the lead's form data in the CRM. No context on who this specific person is, what's likely driving them, what objections to expect, or how to open. The first 10 seconds of the call are generic ("hi, this is X from Scaler") and leads hang up. Both pickup rates and call quality suffer
 4. **The call** — the BDA pulls the lead's details from the CRM and calls them. They try to talk, stretch the call, build connect, understand where the lead is today, sketch a roadmap for where they could be, and answer their questions. Many calls die inside 60 seconds because connect doesn't happen. Others die because the lead asks technical questions about program depth — "*is your AI curriculum updated?*", "*how deep on system design?*" — that the BDA isn't technical enough to answer confidently
 5. **Post-call** — if the call went well, the BDA nudges the lead to take the entrance test. Clearing the test is the gate to the second call and enrolment. **Most leads who had a good call still never take the test** — because nothing between the call and the test builds enough trust to make it worth their time
 6. **Drop-offs happen at every step.** The two places where AI gives the biggest lift are at steps 3 and 5 — prepping the BDA before the call so they open it better, and building lead trust after it so they actually take the test
-

What you'll build

Two pieces. Both delivered over WhatsApp. The pre-sales nudge goes *to the BDA* (no approval step — it's internal). The post-call PDF goes *to the lead*, and every lead-facing send is routed through the BDA for Approve / Edit / Skip — nothing lead-facing fires automatically.

1. Pre-sales WhatsApp nudge to the BDA — ~1 hour

Sent to the BDA on their own WhatsApp, ahead of the call, so they walk in prepared and the first 30 seconds of the call don't sound generic. Better openings drive both pickup rate and call quality up significantly.

Good looks like:

- Short and scannable — the BDA is reading this on their phone 2 minutes before dialling, not studying it
- Specific to *this* lead: who they are in plain English, their likely persona and why, two or three angles that'll resonate (each tied to something real about them), two or three objections to expect with a one-line handle, a suggested opening hook
- Honest about what's inferred vs. fact vs. missing
- Written like a teammate's message, not a corporate memo

No approval gate here — this is the agent briefing the BDA, not speaking to a lead.

2. Post-call personalised PDF, delivered on WhatsApp — ~3 hours

After the call, generate a 2–3 page PDF that:

- Addresses each of the lead's open questions specifically, with evidence (curriculum detail, alumni outcomes from similar learners, concrete ROI reasoning for *their* current package)
- Frames Scaler's strength through the lens of *this* lead's goals, not generic marketing
- Is visibly different lead-to-lead — Rohan's PDF should not look or read like Karthik's
- Is branded, readable, and something a lead would actually open on their phone
- Gets delivered to the lead's WhatsApp with a short personalised covering message

Your agent must accept two input modes:

- **Structured:** a lead profile + a call transcript (text)
- **Audio:** a lead profile + a call recording (audio file)

In the audio path, you transcribe, extract the lead's open questions, and generate the PDF from the extracted content. Whisper, AssemblyAI, Deepgram — your call.

The deployed app must support both input modes. We will upload our own audio/transcript + profile during demo review and see what your agent produces. Don't hardcode for the three standard inputs below.

The three standardised inputs

Use exactly these on your recorded demo so your output is comparable to other candidates'. Each has a profile and a call transcript excerpt — the transcript is what your agent extracts open questions from.

Persona 1 — Rohan Sharma

Profile: Software Engineer, TCS. 4 YoE. Intent: *"want to switch to a product company, tired of service work, interested in AI engineering roles."* LinkedIn: B.Tech CSE VIT Vellore '20, SDE-2 at TCS for 4 years (banking clients: HDFC, Citi), recent AWS Solutions Architect cert.

Call transcript excerpt:

BDA: Rohan, what's bringing you to Scaler? **Rohan:** I've been at TCS for 4 years. Banking projects. I want to move to a product company — and I keep seeing AI engineering roles and wondering if I'm already too late. **BDA:** Not too late. Have you looked at our AI Engineering program? **Rohan:** I've looked. Here's my question though — why should I pay ₹3.5L when Andrew Ng has basically the same stuff for free on Coursera? What's actually different? **BDA:** Good question, let me get back to you on the specifics. **Rohan:** Also — realistically, what salary jump does someone like me get? If I'm going from 14 LPA at TCS to 16 at another service company, the math doesn't work. **BDA:** We have data on that, I'll share. **Rohan:** One more — I want to build real LLM applications. RAG, agents, evals. Is your program on that, or is it more theoretical ML? **BDA:** We'll cover everything you need.

Persona 2 — Karthik Iyer

Profile: Senior Software Engineer, Google. 9 YoE. Intent: *"looking at AI engineering — what's your AI program like for someone like me?"* LinkedIn: IIT Madras CS, 6 years at Google (Search infra), previously Microsoft, frequent open-source contributor.

Call transcript excerpt:

BDA: Karthik, thanks for your time. Tell me what got you interested in Scaler. **Karthik:** Honestly, I'm exploring. I already work at Google. I can read the papers. I just want to make sure I'm not missing anything on the applied side. **BDA:** Of course. What would you want to learn? **Karthik:** My real question is — what would I actually learn here that I can't pick up from papers or internal training? I need to be honest about that before I commit. **BDA:** Our curriculum is very hands-on — **Karthik:** Also — is your cohort going to be at my level? Because if I'm tutoring everyone, I'm not getting value. **BDA:** We have senior folks, yeah. **Karthik:** Last one — are your instructors people who've actually shipped production AI systems, or is it academic? I've sat through enough academic ML.

Persona 3 — Meera Patel

Profile: Final-year B.Tech, Tier-3 college, 0 YoE. Intent: *"need a job, family wants me to take the govt job offer but I want to work at a product company."* LinkedIn: none provided.

Call transcript excerpt:

BDA: Meera, tell me what's on your mind. **Meera:** I'm in my final year. I got a government job offer through campus. My parents want me to take it. But I want to work at a product company. I'm confused. **BDA:** I understand. How can we help? **Meera:** The first thing my parents are going to ask is — can you guarantee I'll get a job after this? Because if I don't, I've turned down a secure government job for

nothing. **BDA:** We have strong placement — **Meera:** And ₹3.5L is more than what my family earns in a year. I genuinely don't know how people afford this. How does that work? **BDA:** We have financing options. **Meera:** Also — I'm nervous about your entrance test. What if I can't clear it? Does that mean I'm not right for this?

Your three PDFs must look and read visibly differently. If they converge to a generic middle, that's a failure signal regardless of how clever the architecture is.

Scaler programs — quick reference

Enough context to sound credible. Don't fabricate details.

- **Scaler Academy** — 12-month flagship, software engineering, 0–7 YoE. ~₹3.5L
- **Scaler Data Science & ML** — 12+ months
- **Scaler DevOps & Cloud** — 9 months

For any curriculum depth, module detail, alumni outcomes, or program specifics — pull from scaler.com. Real grounding beats fabrication. Candidates who scrape or RAG over scaler.com for accurate curriculum claims will stand out on AI craft. A confident wrong answer about curriculum is a hire-blocker — *"let me confirm and get back"* beats fabrication every time.

Demo requirements

Your Loom, under 6 minutes, must show:

1. **Onboarding.** App prompts for evaluator's phone number (we'll give you one) so WhatsApp sends land where we can see them
2. **Pre-sales nudge** generated for each of the three personas and sent to the BDA's WhatsApp (the evaluator's number) — short, scannable, prep material for the call
3. **Post-call PDFs** generated for each of the three personas from their transcripts, side-by-side so personalisation is visibly obvious, BDA approves, real PDFs land on WhatsApp
4. **Audio path** demonstrated once — upload a call recording (you can mock one), show the agent transcribe, extract, generate PDF, deliver
5. **Approval gate** visible on every lead-facing outbound — preview + Approve / Edit / Skip. BDA-facing sends (the pre-sales nudge) don't need the gate

Important for the live evaluation: the deployed app must let us upload our own audio/transcript + persona details and run the full flow. We will test this ourselves. Hardcoding the three standard inputs is a fail.

For WhatsApp with PDF attachments: Twilio WhatsApp Sandbox is the fastest path (~5–10 min including opt-in). Gupshup and MSG91 also work. Fallback: [wa.me](#) link carrying a hosted PDF URL with a visible preview in your UI.

Required in your README

Keep it tight. Three short sections.

1. What you built (plain English, one paragraph).

2. One failure you found (50 words max). Show us an input where your agent produced a bad output, and explain why. We read this section first. An honest failure tells us more than an eval suite would.

3. Scale plan (100 words max). If this went from 1 lead a day to 100,000 leads a month, what breaks first? Pick the one or two real constraints and reason about them. Don't list every possible concern.

One product question (5 minutes, in your submission email)

If you had 30 minutes with us before starting a project at Scaler, what three questions would you ask?

Three questions. No preamble. We're reading for product instinct, not writing skill.

Time and scope

Budget: 5 hours. If you go significantly over, stop and submit what you have with a one-line note.

Non-negotiable for evaluation:

- Pre-sales nudge working across three personas

- Post-call PDF working across three personas, visibly different
- Both input modes (structured + audio) working on the deployed app
- Real PDF lands on WhatsApp during demo
- Deployed link that works for us to test with our own input
- README three sections
- Three questions in the submission email

Cut anything else.

What we look for

Dimension	Weight	What we're reading for
Personalisation depth	30%	Do the three PDFs (and the three nudges) actually look and read differently, or did the agent converge to a generic middle? The strongest single signal
AI craft	25%	Prompt quality, smart model choice, honest handling of uncertainty, active anti-hallucination on curriculum, quality of extraction from transcript/audio
Shipping velocity	15%	Deployed, working end-to-end in 5 hours. PDF actually lands on WhatsApp. App works when we test it with our own input
Product thinking	15%	Your three questions. Your 50-word failure. Did you understand the real job is building enough trust for the lead to take the test?

Dimension	Weight	What we're reading for
AI nativeness	10%	Fluent use of current tooling — frontier models, STT, PDF generation, WhatsApp APIs, AI dev environments (Cursor, Claude Code, v0). Signal shows up in how much you shipped in 5 hours
Communication	5%	Tight README, clean Loom, decisions documented

What would disqualify

- No live deployed link, or a deployed link that breaks when we test with our own input
 - PDF doesn't actually reach the evaluator's WhatsApp in the demo
 - Three persona PDFs that look visibly interchangeable
 - Confident hallucinated claims about Scaler curriculum
 - Lead-facing messages firing without a BDA approval step
 - Audio path not actually implemented (we will test it)
 - A README that's clearly LLM-generated and not reviewed
-

How to submit

Email **both**:

- amar.srivastava@scaler.com
- abhimanyu@scaler.com

Subject: [AI builder](#)

Body of the email contains:

1. Your name, LinkedIn, one line on why you want this role
2. Your **three questions** (from the product question above)
3. Live deployed link
4. GitHub repo link (public, or private with access to both emails above)

5. Loom under 6 minutes
6. README link (in the repo)

If your submission email runs longer than a screen of text, you're probably doing it wrong.

Honest notes

- 5 hours means 5 hours. Respect for constraints is part of the role
- Use AI tools aggressively. 90% prompting, 10% fixing is the ratio we want to see
- Where something is ambiguous, decide and document

Good luck.